

# Operating System

## Module 2. Process Concept and scheduling

**Dr. Prasanna Shete**

Dept. of Computer Engineering

[prasannashete@somaiya.edu](mailto:prasannashete@somaiya.edu)

Mobile/WhatsApp 9960452937



# Processes and Threads

- Processes have two characteristics:
- **Resource ownership**– ownership of resources like main memory, I/O channels/devices and files
  - process includes a virtual address space to hold the process image
- **Scheduling/execution**- Execution of process follows an execution path (trace) through one or more programs;
  - may be interleaved with other processes
- Process has execution state and dispatching priority
- These two characteristics are treated independently by the operating system

# Processes and Threads

- The unit of dispatching is referred to as a ***thread*** (or lightweight process)
- The unit of resource ownership is referred to as a ***process*** (or ***task***)

# Multithreading

- The ability of an OS to support multiple, concurrent paths of execution **within a single process**.
- 4 approaches (relationship between threads and processes)
  - 1 Process, 1 Thread (e.g MS-DOS)
  - 1 Process, Multiple Threads (JRE)
  - Multiple Processes, 1 Thread / Process (Unix)
  - Multiple Processes, Multiple Threads / Process (Linux, Windows, Solaris)

# Multithreading

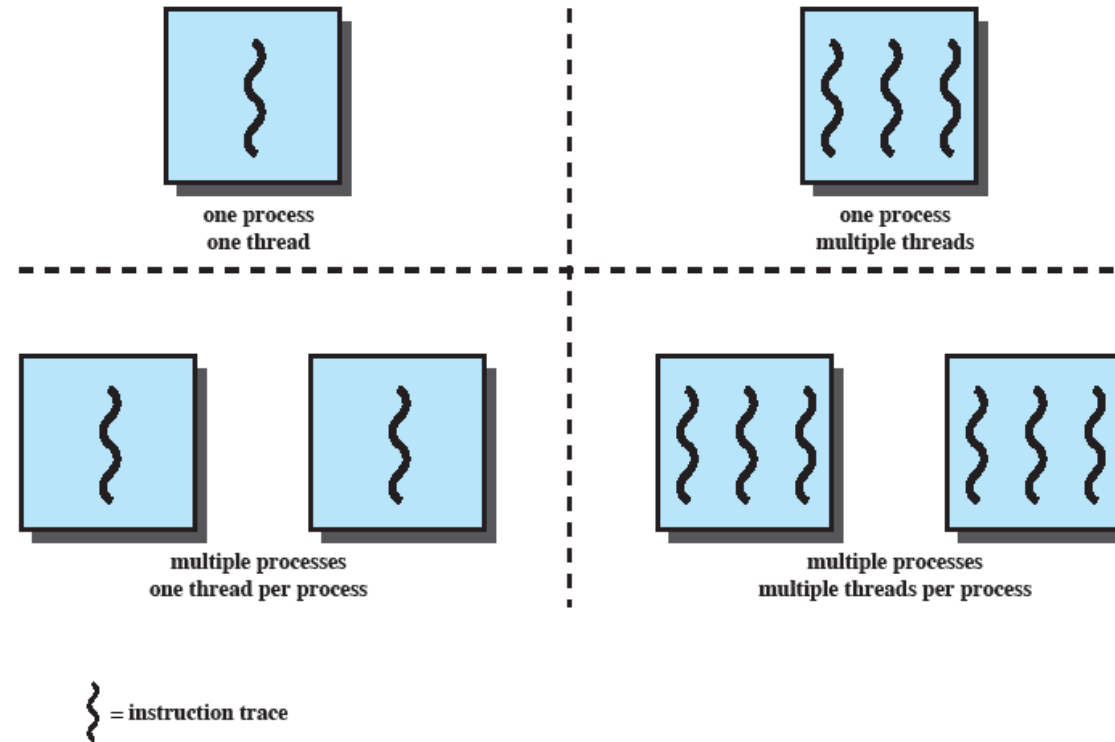


Figure 4.1 Threads and Processes [ANDE97]

# Multithreading

- In multithreaded environment a **Process is defined as the unit of resource allocation and a unit of protection;**
  - following things are associated with it:
- A virtual address space that holds the process image
- Protected access to
  - Processors
  - Other processes
  - Files
  - I/O resources

# Threads within a process

- Each thread has
  - An execution state (running, ready, etc.)
  - Saved thread context when not running  
(can be viewed as independent PC operating within a process)
  - An execution stack
  - Some per-thread static storage for local variables
  - Access to the memory and resources of its process (shared with all threads in that process)

# Threads V/s processes

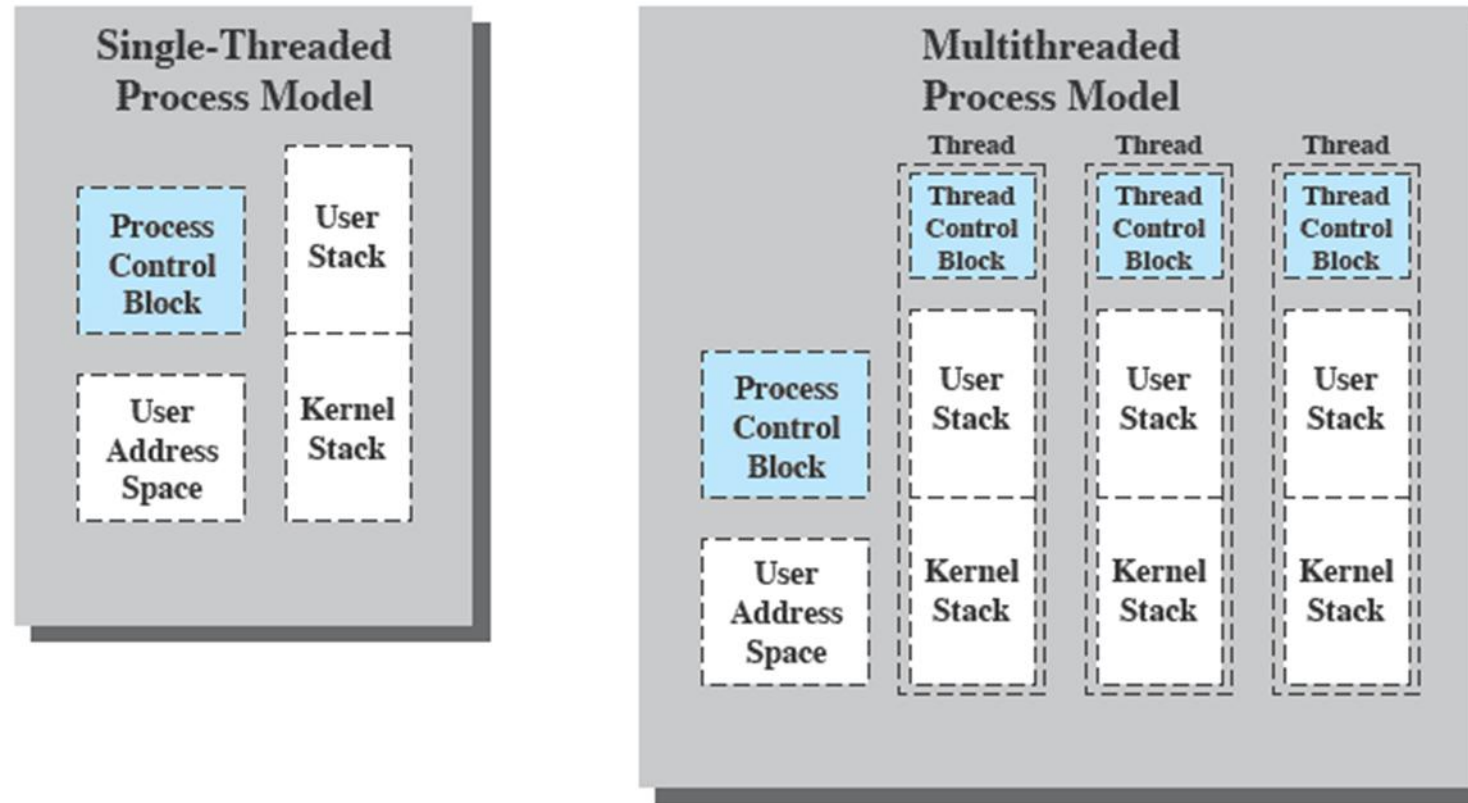


Figure 4.2 Single Threaded and Multithreaded Process Models



# Distinction between threads and processes from the point of view of process management

- **Single-threaded process model-**
- No distinct concept of thread
- the representation of a process includes
  - its PCB (process control block)
  - user address space,
  - user and kernel stacks to manage the call/return behaviour of the execution of the process.

While the process is running, it controls the processor registers

When the process is not running, Contents of these registers are saved

# Distinction between threads and processes from the point of view of process management

- **In a multithreaded environment** there is single PCB and user address space associated with the process, **but** separate stacks for each thread
- Separate control block called **Thread Control Block**, is maintained for each thread containing **register values, priority, and other thread-related state information**

# Distinction between threads and processes from the point of view of process management

- **Multithreaded environment...**
  - All threads of a process share the state and resources of that process
  - They reside in the same address space and have access to the same data.
  - When one thread alters data in the memory, other threads see the results, as and when they access that location
  - Threads in the same process can read the contents of the same memory location concurrently
    - If one thread opens a file with read privileges, other threads in the same process can also read from that file

# Benefits to performance

- **Less creation time-** Takes less time to create a new thread than a process (10 times faster)
- **Less termination time-** time required to terminate a thread is less than that for process
- **Less Switching time-** Switching between two threads within the same process takes less time than switching processes
- **Efficiency enhancement in communication between different executing programs**
  - Threads within a process share memory and files, thus can communicate with each other without invoking the kernel

# Example scenarios of Threads in Single-User System

- Foreground and background work
  - Spreadsheet: One thread displays menu, other executes user commands
- Asynchronous processing
  - Power failure protection; write contents of RAM to disk every 100ms
- Speed of execution
  - Compute one batch of data while reading next batch; multi-processor scenario
- Modular program structure

# Thread Functionality

- Thread States- Running, Ready, Blocked; Suspended state is not associated with threads
  - If a process is suspended, all its threads are suspended
- Thread operations:
  - Spawn (another thread)
  - Block
    - Issue: will blocking a thread block other, or *all*, threads
  - Unblock
  - Finish (thread)
    - Deallocate register context and stacks

# Multithreading on a Uniprocessor

- Multithreading enables interleaving of multiple threads within multiple processes

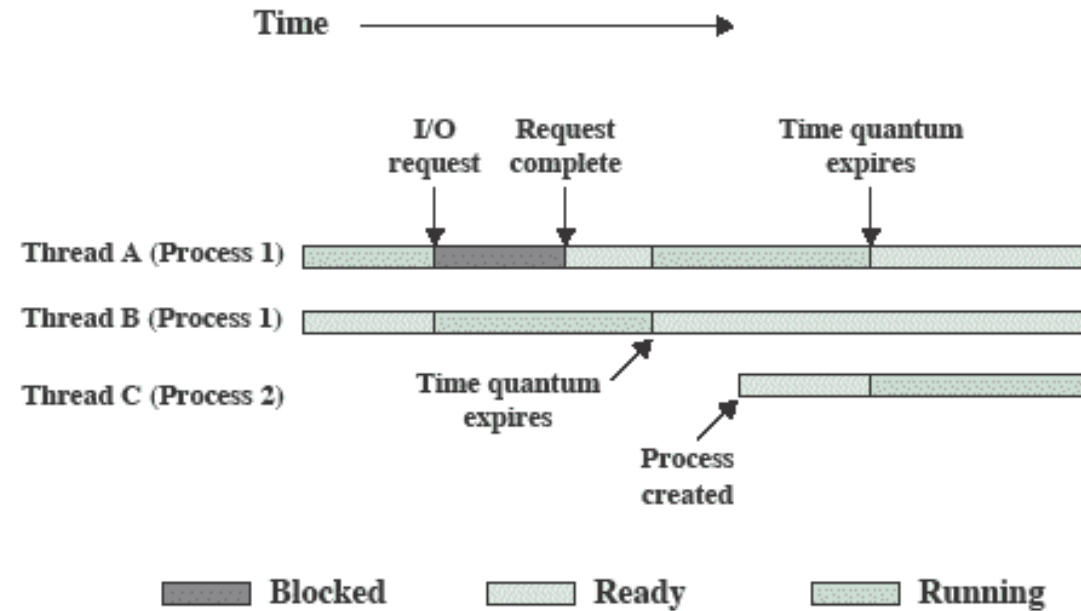


Figure 4.4 Multithreading Example on a Uniprocessor


# Thread Implementation: Categories

- User Level Thread (ULT)
- Kernel level Thread (KLT) also called:
  - kernel-supported threads
  - lightweight processes



# Thread Implementation: ULTs

- All thread management done by the application
  - The kernel is not aware of the existence of threads
  - Threads Library: package of application level routines for thread management
- Kernel schedules process as unit



(a) Pure user-level

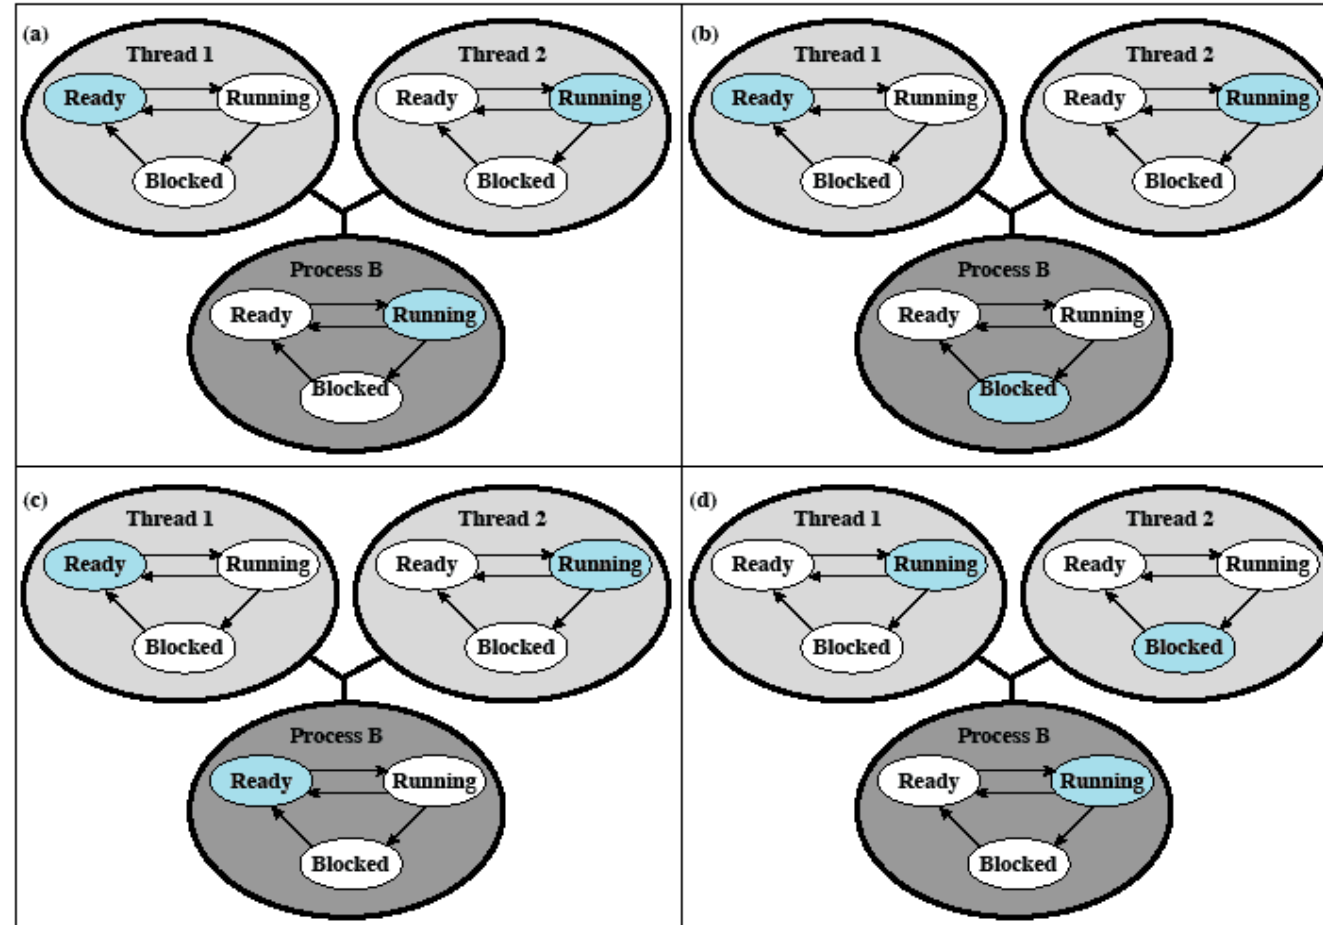
# ULTs

- Threads Library: package of application level routines for thread mngmentsin
  - contains the code for creating and destroying threads, for passing messages and databetween threads, for scheduling threads and for saving and restoring thread contexts
- Application begins with single thread and starts running in that thread
- The application and all its threads are allocated to a **single process** managed by the kernel
- When the application is running it may spawn a new thread to run within the same process

# ULTs

- Threads Spawing- is carried out by invoking the spawn utility in the threads library; Control passed to that utility by a procedure call
- Threads library creates the data structure for the new thread and then passes control to one of the threads that is in the Ready state using some scheduling algorithm
- When control is passed to the library, the context of current thread is saved; and restored when control is passed back to the thread
- These activities take place in the user space; kernel is unaware
- Kernel continues to schedule the process as a unit

# Relationships between ULT Thread and Process States



Colored state  
is current state

Figure 4.7 Examples of the Relationships Between User-Level Thread States and Process States

# Relationships between ULT Thread and Process States

- Suppose Process B has 2 threads and is executing in Thread 2
  1. Process B blocked: During execution the application makes a system call that blocks process B → kernel switches control to another process
    - Thread 2 will still be in the running state- not actually executing on the processor- but perceived as being in the running state by the threads library
  2. Clock interrupt → process B moved to ready state by kernel; switches control to another process
    - Data structures maintained by threads library indicate that thread 2 is in running state
  3. Thread blocked: Thread 2 needs some action to be performed by thread 1, → Thread 2 is blocked; thread 1 transitions from Ready to running state
    - Data structures maintained by threads library are updated

# ULT Advantages


- Thread switching does not need Kernel level privileges
  - Thread management data structures are within user space of process
- Scheduling can be application specific
  - Scheduling algorithms can be tailored to the application without disturbing underlying OS scheduler
- ULTs can run on any OS
  - No changes in the kernel are required to support ULT

# ULT Disadvantages

- Multi threaded applications cannot take advantage of multiprocessing
- When one thread **blocks on a system call**, all threads within that process are blocked
- Remedy:
- Write application with multiple processes rather than threads
  - → eliminates the advantage of threads
- **Jacketing**- Converts a blocking system call into non-blocking system call
  - Rather than directly calling the system I/O routine, thread calls an application level I/O jacket routine; checks the availability of I/O

# Thread Implementation: KLTs

- Kernel maintains context information for the process and the threads
  - No thread management done by application
- Scheduling is done on a thread basis
  - Windows is an example of this approach



(b) Pure kernel-level




# Advantages of KLT

- Kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread in a process is blocked, the kernel can schedule another thread of the same process
- Kernel routines themselves can be multithreaded
- Disadvantages
- The transfer of control from one thread to another within the same process requires a mode switch to the kernel

# Thread Implementation: Combined Approach

- Thread creation done in the user space
- Bulk of scheduling and synchronization of threads by the application
- Multiple ULTs from single application are mapped onto smaller (or equal ) number of KLTs
- Multiple threads within same application can run in parallel on multiple processors
  - Blocking system call need not block the entire process
- Example: Solaris



(c) Combined

# Relationship Between Thread and Processes

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX

# Categories of Computer Systems

- Single Instruction Single Data (SISD) stream
  - Single processor executes a single instruction stream to operate on data stored in a single memory
- Single Instruction Multiple Data (SIMD) stream
  - Each instruction is executed on a **different set of data** by different processors
- Multiple Instruction Single Data (MISD) stream
  - A sequence of data is transmitted to a set of processors, each of them execute a different instruction sequence
- Multiple Instruction Multiple Data (MIMD)
  - A set of processors simultaneously execute **different instruction sequences on different data sets**

# Categories of Computer Systems

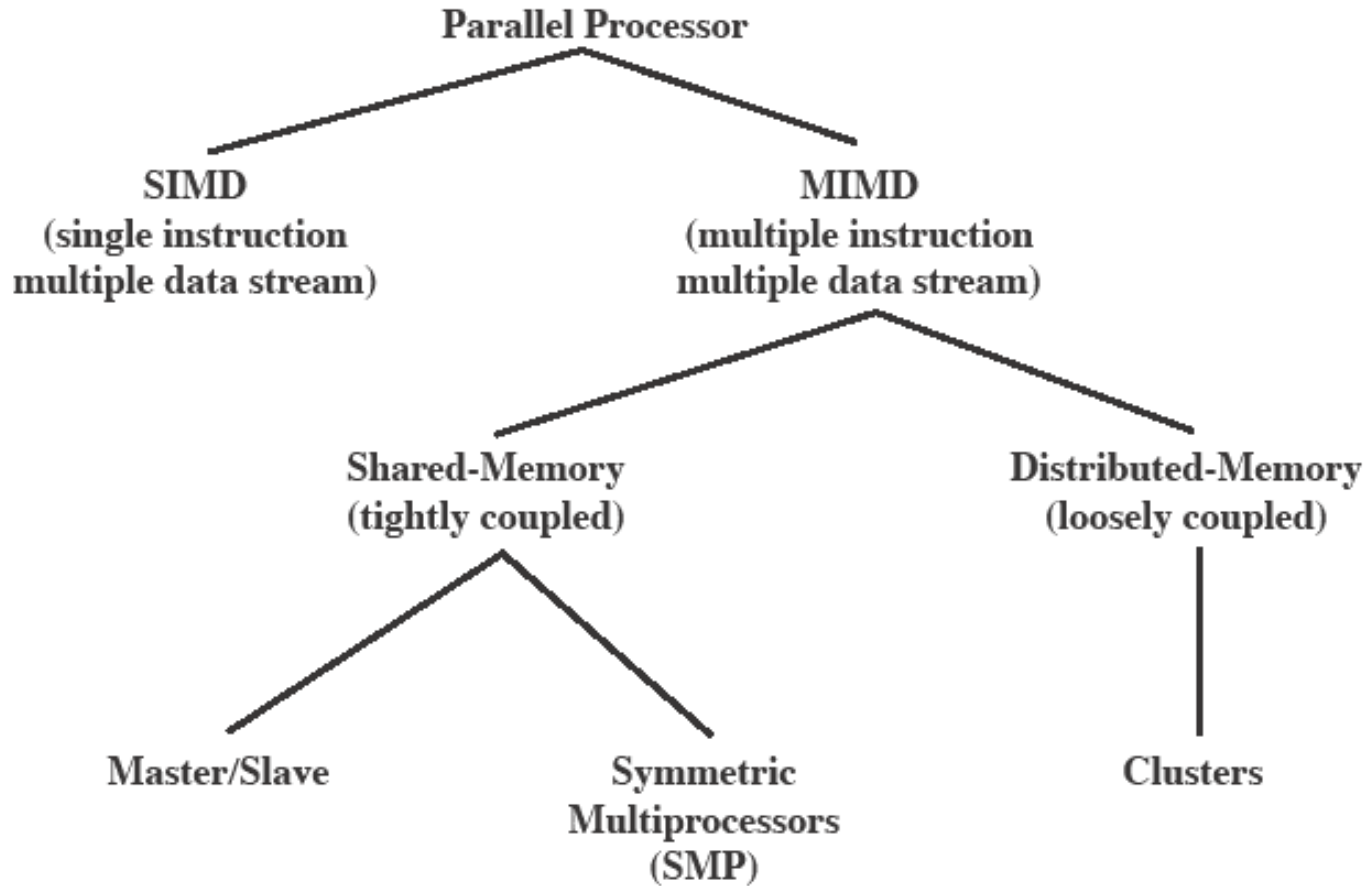
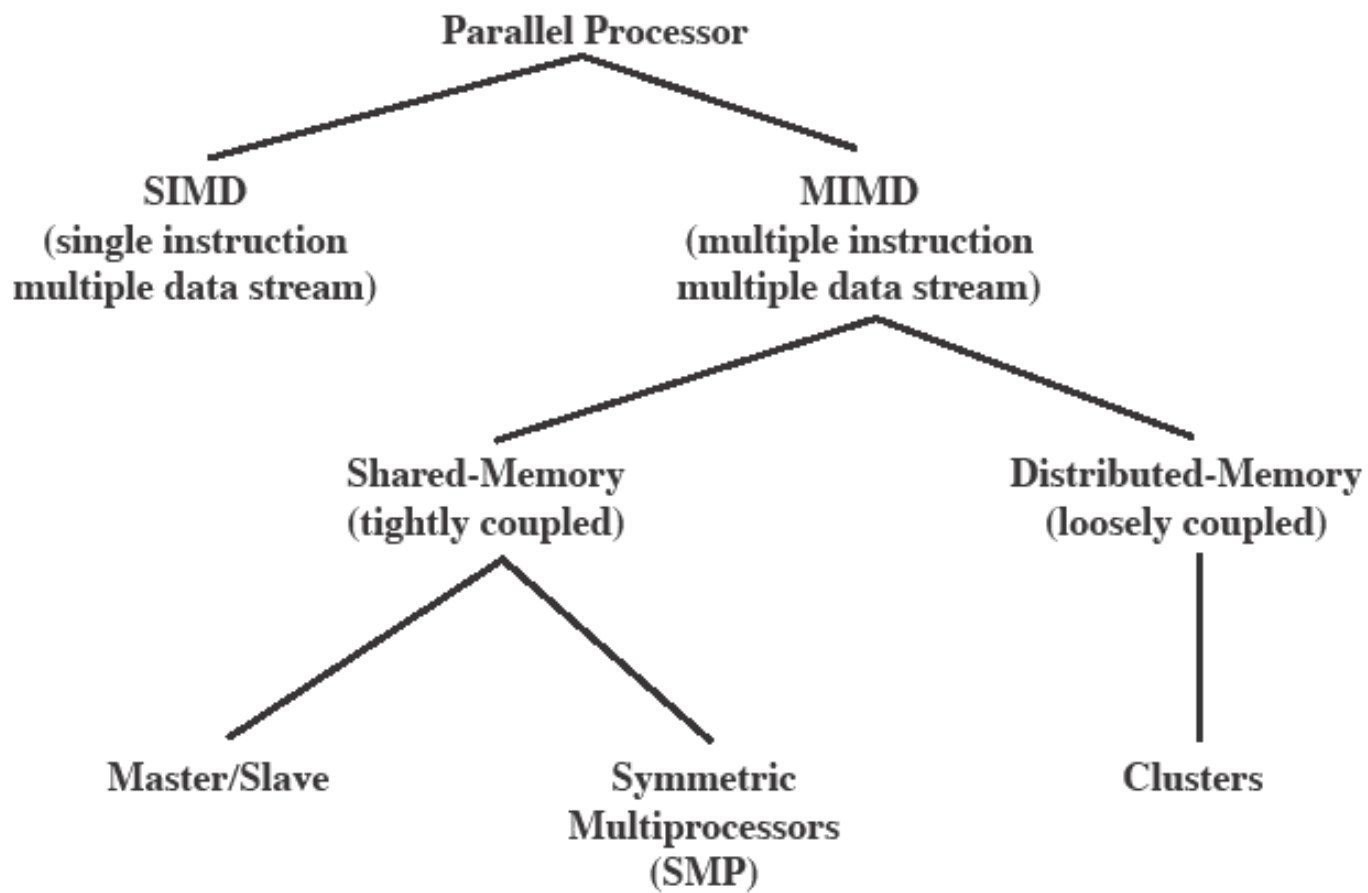


Figure 4.8 Parallel Processor Architectures



**Figure 4.8 Parallel Processor Architectures**

# Symmetric Multiprocessing

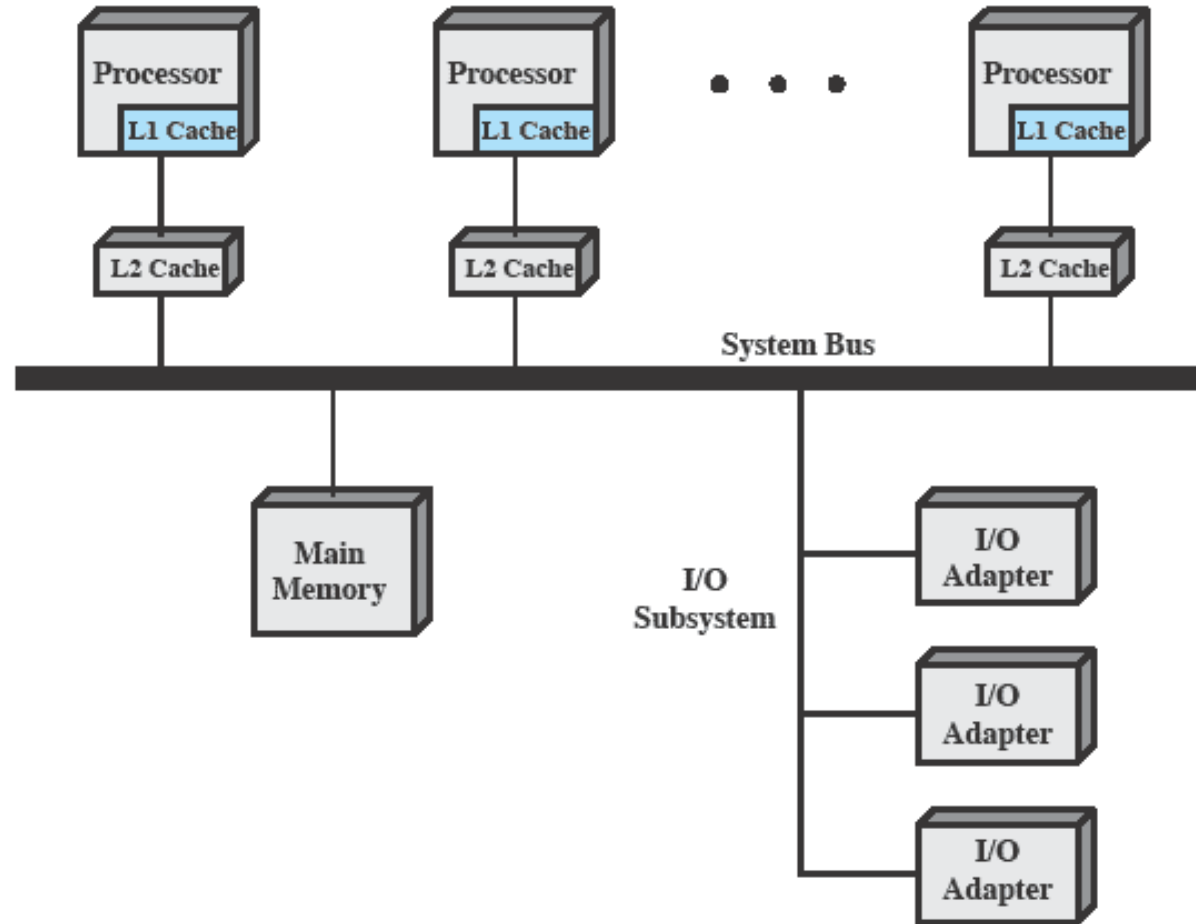


Figure 4.9 Symmetric Multiprocessor Organization

# Symmetric Multiprocessing

- In a symmetric multiprocessor, Kernel can execute on any processor
  - Each processor does **self-scheduling** from the pool of available process or threads (typically)
- The kernel can be constructed as multiple processes or multiple threads, allowing portions of the kernel to execute in parallel
- Complicates the OS
  - must **ensure that two processors do not choose the same process** and that processes are not somehow lost from the queue.
  - Techniques must be employed to resolve and synchronize claims to resources



# Multiprocessor OS Design Considerations

- The key design issues include
  - Simultaneous concurrent processes or threads
  - Scheduling
  - Synchronization
  - Memory Management
  - Reliability and Fault Tolerance

# Multiprocessor OS Design Considerations

## 1. Simultaneous concurrent processes or threads:

- Kernel routines need to be re-entrant- to allow several processors to execute the same kernel code simultaneously
- With multiple processors executing the same or different parts of the kernel, kernel tables and management structures must be managed properly to avoid deadlock or invalid operations

## 2. Scheduling:

- Scheduling may be performed by any processor, so conflicts must be avoided
- If kernel-level multithreading is used, then there is opportunity to schedule multiple threads from the same process simultaneously on multiple processors

# Multiprocessor OS Design Considerations

## 3. Synchronization:

- With multiple active processes having potential access to shared address spaces or shared I/O resources, care must be taken to provide effective synchronization
- Synchronization is a facility that enforces mutual exclusion and event ordering
  - A common synchronization mechanism used in multiprocessor operating systems is locks

# Multiprocessor OS Design Considerations

## 4. Memory management:

- Must deal with all of the issues found on uniprocessor computers
- OS needs to exploit the available hardware parallelism, such as multiported memories, to achieve the best performance
- The paging mechanisms on different processors must be coordinated to enforce consistency when several processors share a page or segment and to decide on page replacement

## 5. Reliability and fault tolerance:

- OS should provide **graceful degradation in the case of processor failure**
- The scheduler and other portions of the OS must recognize the loss of a processor and restructure management tables accordingly

# Multiprocessor Scheduling Design Issues

- Scheduling on a multiprocessor involves 3 interrelated issues:
  - Assignment of processes to processors
  - Use of multiprogramming on individual processors
  - Actual dispatching of a process
- Approach taken depends on the degree of granularity of applications and the number of processors available

# Assignment of Processes to Processors

- Assuming all processors are equal, treat processors as a pooled resource and
- assign process to processors on demand
  - Should the assignment be static or dynamic?
- Dynamic Assignment:
  - threads are moved from a queue for one processor to a queue for another processor

# Assignment of Processes to Processors

- Static Assignment:
- Permanently assign of process/thread to a processor
  - Dedicated short-term queue for each processor
  - Less overhead
  - Allows the use of 'group' or 'gang' scheduling (see later)
- May result in processor remaining idle; while others have a backlog/ pending work
  - Solution: use a common queue
  - All processes go into one global queue and are scheduled on to any available processor

# Assignment of Processes to Processors

- Both dynamic and static methods require some way of assigning a process to a processor
- Two methods:
  - Master/Slave
  - Peer-to-peer
- There are a spectrum of approaches between these two extremes



# Assignment of Processes to Processors:

- **Master / Slave Architecture**

- Key kernel functions always run on a particular processor- Master
- Master is responsible for scheduling
- Slave sends service request to the master

## Disadvantages

- Failure of master brings down whole system
- Master can become a performance bottleneck

- **Peer Architecture:**

- Kernel can execute on any processor
- Each processor does self-scheduling
- Complicates the operating system
  - Need to make sure two processors do not choose the same process

# Process Scheduling

- Usually processes are not dedicated to processors
- A single queue is used for all processes
- Or multiple queues are used for differing priorities
  - All queues feed to the common pool of processors

# Thread Scheduling

- Threads execute separate from the rest of the process
- An application can be a set of threads that cooperate and execute concurrently in the same address space
  - Dramatic gains in performance in multi-processor systems, compared to running in uniprocessor system
- Four general approaches:
  - Load Sharing
  - Gang Scheduling
  - Dedicated Processor assignment
  - Dynamic scheduling

# Thread Scheduling: Load Sharing

- Threads are not assigned to a particular Processor
- Global queue of ready threads is maintained
- Each processor, when idle, selects a thread from the queue
- Advantages:
  - Simple Approach
  - Load is distributed evenly across the processors
  - No centralized scheduler required
    - whenever a processor is available, scheduling routine of OS runs on that processor to select the next thread
  - The global queue (of tasks) can be organized and accessed using any scheduling algorithm

# Thread Scheduling: Load Sharing

- 3 different versions of load sharing approach

## – FCFS

- Each thread of a job is placed consecutively at the end of shared queue
- When processor becomes idle, it picks up next ready thread and executes it till completion (or blocking)

## – Smallest Number of Threads First

- Priority based shared queue; highest priority to threads of that job which has least no. of threads

## – Pre-emptive Smallest Number of Threads First

- Job with smallest no. of threads than the job which is already in execution, will pre-empt the threads of scheduled job

# Thread Scheduling: Load Sharing

- **Disadvantages:**
- Central queue resides in the region of memory that needs to be accessed in a mutually exclusive manner
  - Can lead to bottlenecks
- Preemptive threads are unlikely to resume execution on the same processor
  - Caching becomes less efficient
- If all threads are in the global queue, all threads of a given program will not gain access to the processors at the same time

# Thread Scheduling: Gang Scheduling

- Scheduling set of threads simultaneously on a set of processors
  - A set of related threads/processes is scheduled to run on a set of processors at the same time
- Parallel execution of closely related processes may reduce overhead such as: process switching and synchronization blocking
  - Improves Performance
- Gang scheduling is usually applied to simultaneous scheduling of threads that make up a single process
  - Used for medium to fine grained parallel applications

# Thread Scheduling: Gang Scheduling

- Involves Following activities:
  1. Group of related threads are scheduled as a unit (or gang)
  2. All members of gang run simultaneously on different processors
  3. All members start and end their time slices together
- Gang scheduling works because all CPUs are scheduled synchronously, by dividing time into discrete quanta



# Thread Scheduling: Dedicated Processor Assignment

- Opposite to load sharing
- Implicit scheduling defined by Assignment of threads to processors
- Each program is allocated with
- Number of Processors = No. of threads in the program for the duration of program execution
  - When program terminates processors return to general pool; these processors can be allocated to other programs
- Some processors may remain idle
- No multiprogramming of processors

# Thread Scheduling: **Dynamic Scheduling**

- Number of threads in a process may be/are altered dynamically by the application during the course of execution
  - By using some language/ system tools; e.g a thread giving birth to many child threads
- Dynamic Scheduling allows the OS to adjust the load to improve utilization/performance
- OS as well as application are involved in scheduling decisions



