



Course Name:	Data Analysis Laboratory (216H03L501)	Semester:	V
Date of Performance:	13/10/2025	DIV/ Batch No:	HDA2
Student Name:	Aaryan Sharma	Roll No:	16010123012

TITLE: Perform a time series analysis on health care (AR/MA/ARIMA)

AIM: To perform forecasting using time series analysis

Expected OUTCOME of Experiment:

CO4: Perform Time series Analytics and forecasting

Books/ Journals/ Websites referred:

<https://www.geeksforgeeks.org/machine-learning/time-series-analysis-and-forecasting/>

Pre Lab/ Prior Concepts:

Students should have a basic understanding of: Time series Analytics and forecasting

Procedure:

Data set Used: Hospital_patients_datasets

Step1: Select and Load the dataset

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from prophet import Prophet

# Load the dataset
file_path = 'Hospital_patients_datasets.csv' # replace with actual path
data = pd.read_csv(file_path)

/tmp/ipython-input-1011798746.py:8: DtypeWarning: Columns (0) have mixed types. Specify dtype option on import or set low_memory=False.
data = pd.read_csv(file_path)
```

Step 2: Convert 'ScheduledDay' and 'AppointmentDay' to datetime format

```
# Convert 'ScheduledDay' and 'AppointmentDay' to datetime format
data['ScheduledDay'] = pd.to_datetime(data['ScheduledDay'])
data['AppointmentDay'] = pd.to_datetime(data['AppointmentDay'])

# Create new features: Appointment Weekday, Scheduled Weekday, Days Between Scheduling and Appointment
data['AppointmentWeekday'] = data['AppointmentDay'].dt.day_name()
data['ScheduledWeekday'] = data['ScheduledDay'].dt.day_name()
data['DaysBetween'] = (data['AppointmentDay'] - data['ScheduledDay']).dt.days
```

Step 3: Forecasting Daily Attendance

```
# Forecasting Daily Attendance
# Step 1: Preprocess the data to aggregate daily attendance
data['Attended'] = data['No-show'].apply(lambda x: 0 if x == 'Yes' else 1)
attendance_daily = data.groupby('AppointmentDay')['Attended'].sum().reset_index()
# Step 2: Rename columns for Prophet ('ds' for date, 'y' for the target)
attendance_daily.columns = ['ds', 'y']
# Ensure the 'ds' column is in datetime format and without timezone
attendance_daily['ds'] = pd.to_datetime(attendance_daily['ds']).dt.tz_localize(None)
# Ensure 'y' is numeric
attendance_daily['y'] = pd.to_numeric(attendance_daily['y'], errors='coerce')
# Drop any rows with missing values (NaNs)
attendance_daily = attendance_daily.dropna()
```



Step 4: Initialize Prophet model for forecasting

```
# Step 3: Initialize Prophet model for forecasting
model = Prophet()
```

Step 5: Fit the model

```
# Step 4: Fit the model
model.fit(attendance_daily)

INFO:prophet:Disabling yearly seasonality. Run prophet with yearly_seasonality=True to override this.
INFO:prophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.
INFO:cmdstanpy:cmdstanpy:running CmdStan, num_threads: None
INFO:cmdstanpy:cmdstanpy:args: ['~/usr/local/lib/python3.12/dist-packages/prophet/stan_model/prophet_model.bin', 'random', 'seed-77622', 'data', 'file~/tmp/tmpafqr88a/gjb5if9x.json', 'init~/tmp/tmpafqr88a/1654d5en.json', 'output
09:23:11 - cmdstanpy - INFO - Chain [1] start processing
09:23:11 - cmdstanpy - INFO - Chain [1] done processing
INFO:cmdstanpy:Chain [1] done processing
INFO:cmdstanpy:Chain [1] done processing
prophet.forecaster.prophet at 0x907d8592810]
```

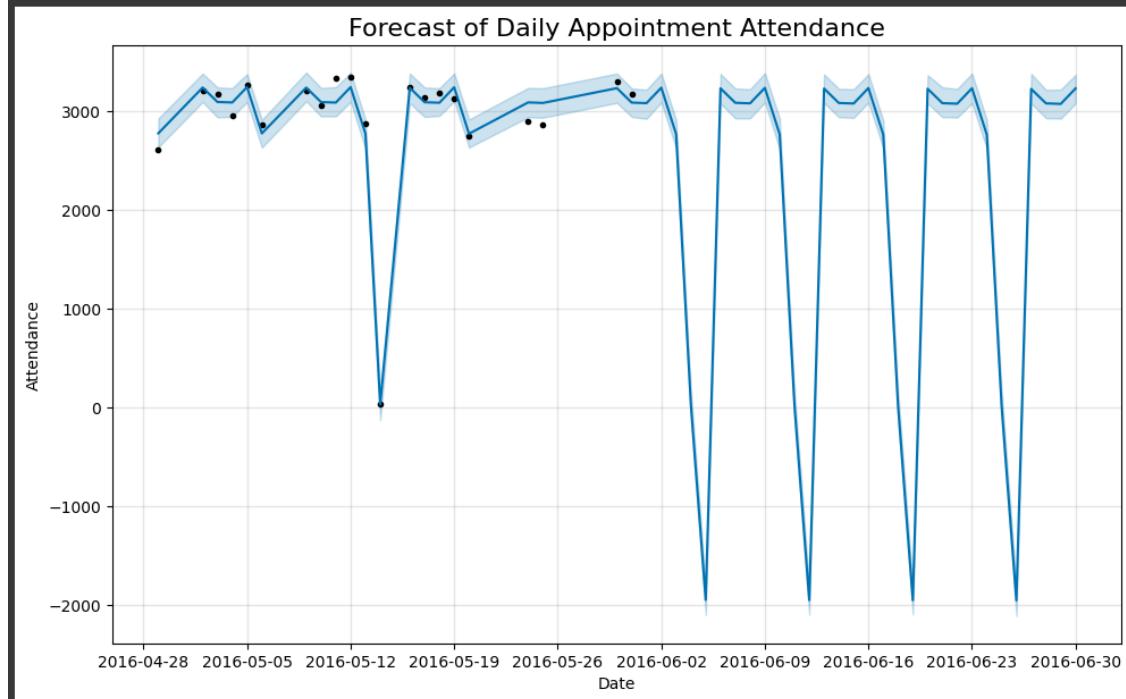
Step 6: Predict future attendance

```
# Step 5: Create future dates for prediction (e.g., next 30 days)
future_dates = model.make_future_dataframe(periods=30)

# Step 6: Predict future attendance
forecast = model.predict(future_dates)
```

Step 7: Plot the forecast

```
# Step 7: Plot the forecast
fig = model.plot(forecast)
plt.title('Forecast of Daily Appointment Attendance', fontsize=16)
plt.xlabel('Date')
plt.ylabel('Attendance')
plt.show()
```





Step 8: Exploratory Data Analysis Functions

```
# Exploratory Data Analysis Functions
def plot_no_show_distribution():
    """Plot no-show distribution"""
    plt.figure(figsize=(8, 6))
    sns.countplot(x='No-show', data=data, palette="Set2")
    plt.title('Distribution of No-shows', fontsize=16)
    plt.ylabel('Count')
    plt.xlabel('No-show Status')
    plt.show()

def plot_age_distribution():
    """Plot age distribution of patients"""
    plt.figure(figsize=(10, 6))
    sns.histplot(data=data, x='Age', bins=50, kde=True, color='skyblue')
    plt.title('Age Distribution of Patients', fontsize=16)
    plt.xlabel('Age')
    plt.ylabel('Count')
    plt.show()

def plot_age_vs_no_show():
    """Plot relationship between Age and No-show status"""
    plt.figure(figsize=(10, 6))
    sns.boxplot(x='No-show', y='Age', data=data, palette="Set3")
    plt.title('Age vs. No-show Status', fontsize=16)
    plt.xlabel('No-show Status')
    plt.ylabel('Age')
    plt.show()

def plot_medical_conditions_vs_no_show():
    """Plot relationship between medical conditions and No-show status"""
    plt.figure(figsize=(12, 6))
    medical_conditions = ['Hipertension', 'Diabetes', 'Alcoholism', 'Handcap']

    for condition in medical_conditions:
        plt.figure()
        sns.countplot(x='No-show', hue=condition, data=data, palette="Set2")
        plt.title(f'Relationship between {condition} and No-show')
        plt.ylabel('Count')
        plt.xlabel('No-show Status')
        plt.legend(title=condition)
        plt.show()

def plot_sms_vs_no_show():
    """Plot relationship between SMS reminders and No-show status"""
    plt.figure(figsize=(10, 6))
    sns.countplot(x='No-show', hue='SMS_received', data=data, palette="Set2")
    plt.title('Relationship between SMS Received and No-show', fontsize=16)
    plt.ylabel('Count')
    plt.xlabel('No-show Status')
    plt.legend(title='SMS Received', loc='upper right')
    plt.show()
```

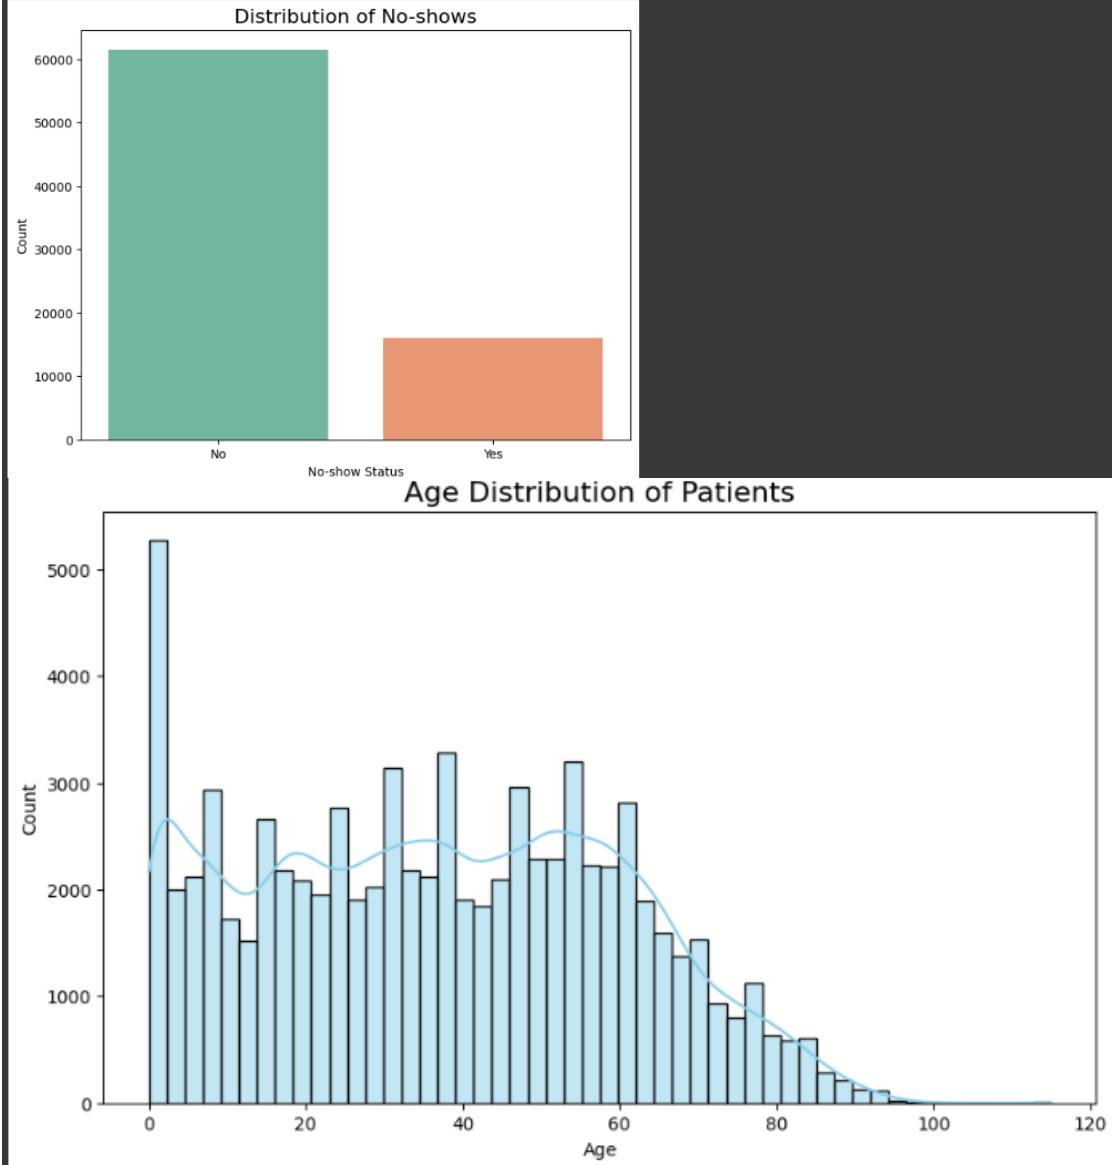


Step 9: Running the analysis functions

```
# Running the analysis functions
plot_no_show_distribution()
plot_age_distribution()
plot_age_vs_no_show()
plot_medical_conditions_vs_no_show()
plot_sms_vs_no_show()

/tmp/ipython-input-752034224.py:5: FutureWarning:
Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.

sns.countplot(x='No-show', data=data, palette="Set2")
```



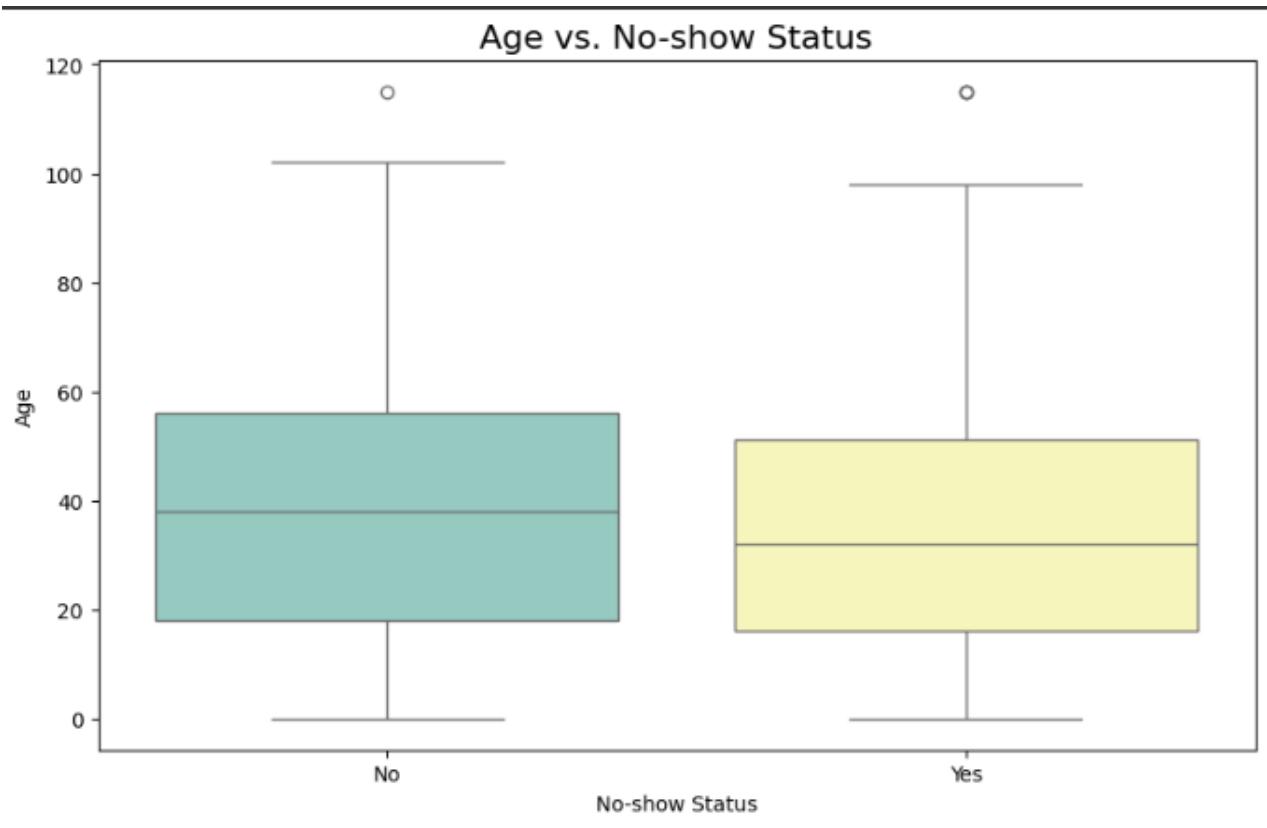


SOMAIYA
VIDYAVIHAR UNIVERSITY

K J Somaiya School of Engineering
(formerly K J Somaiya College of Engineering)

K. J. Somaiya School of Engineering, Mumbai-77
(Somaiya Vidyavihar University)
Department of Computer Engineering

Somaiya
T R U S T



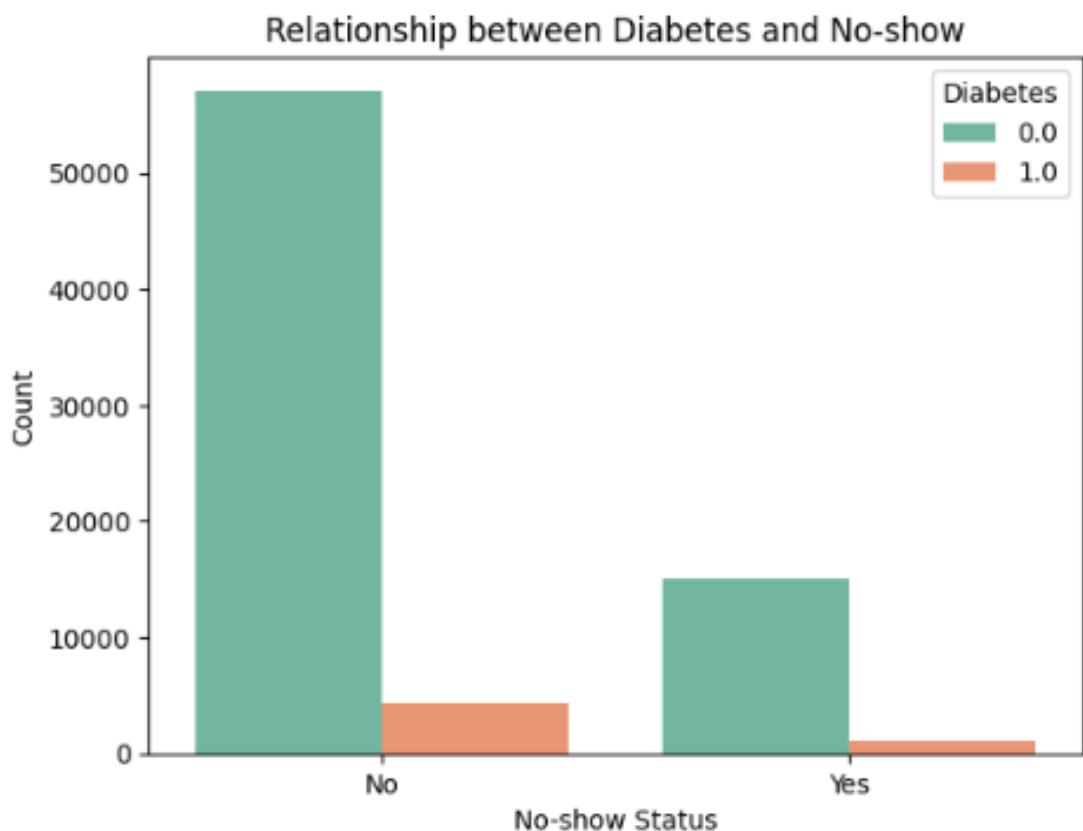
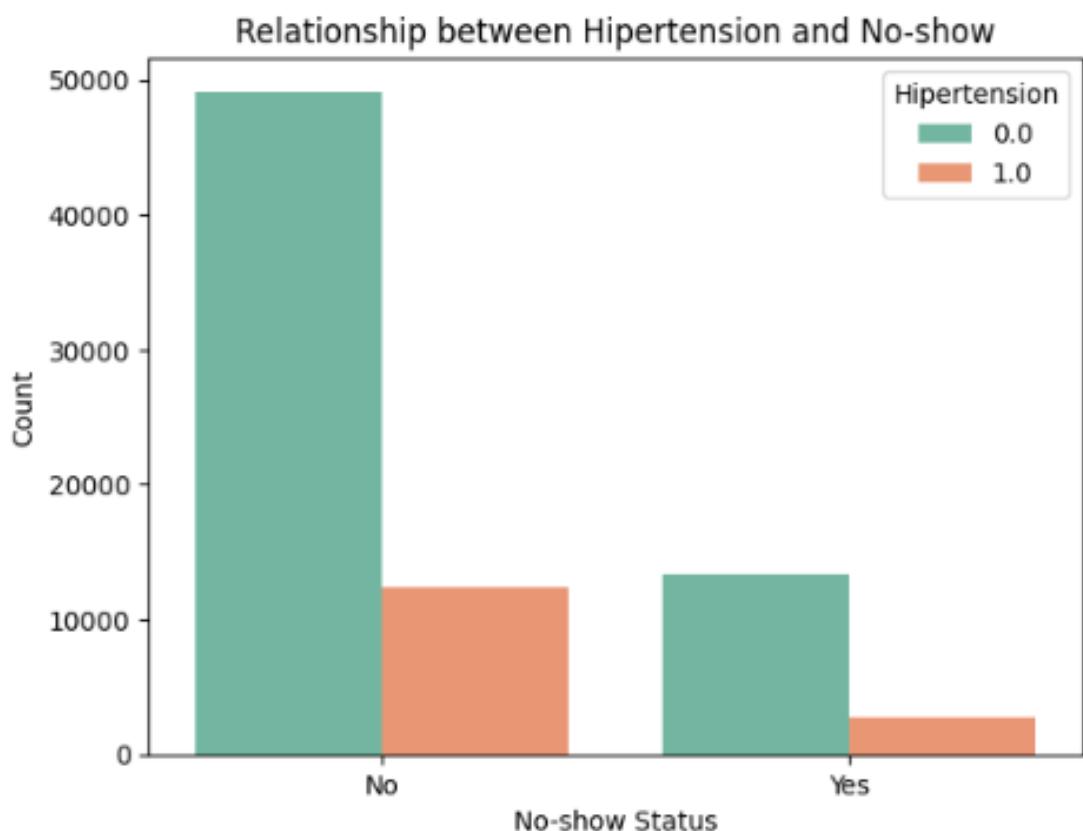


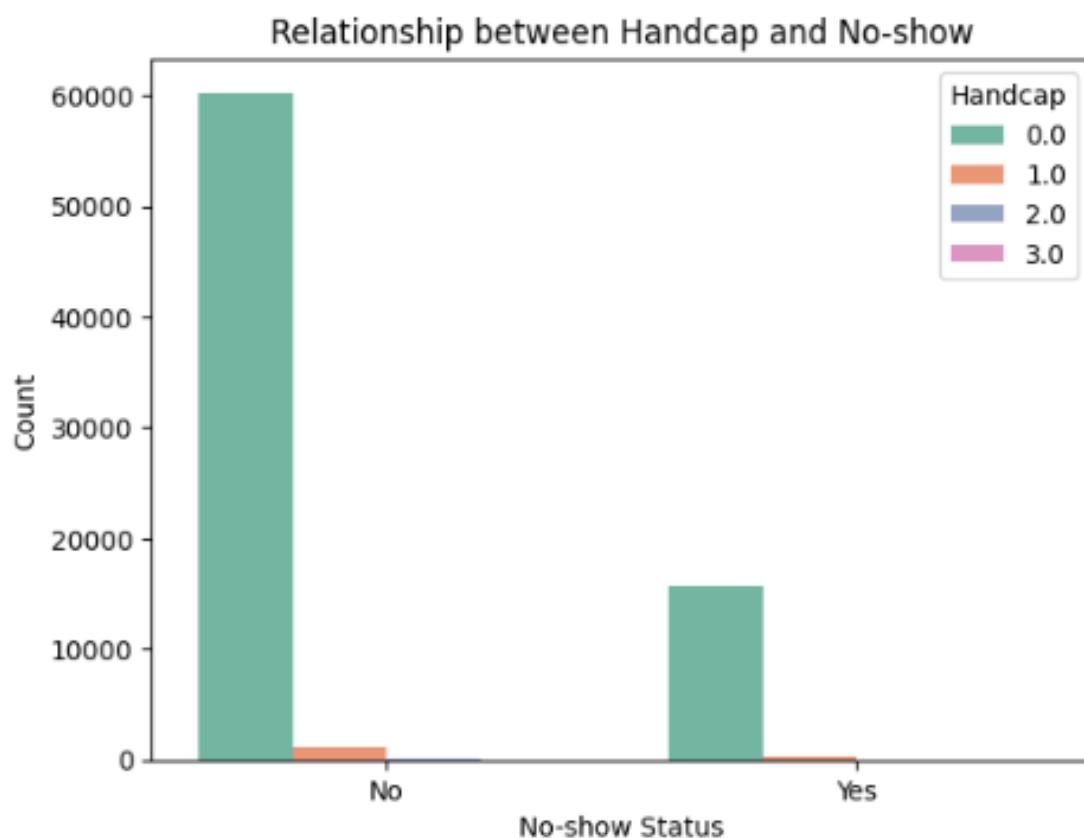
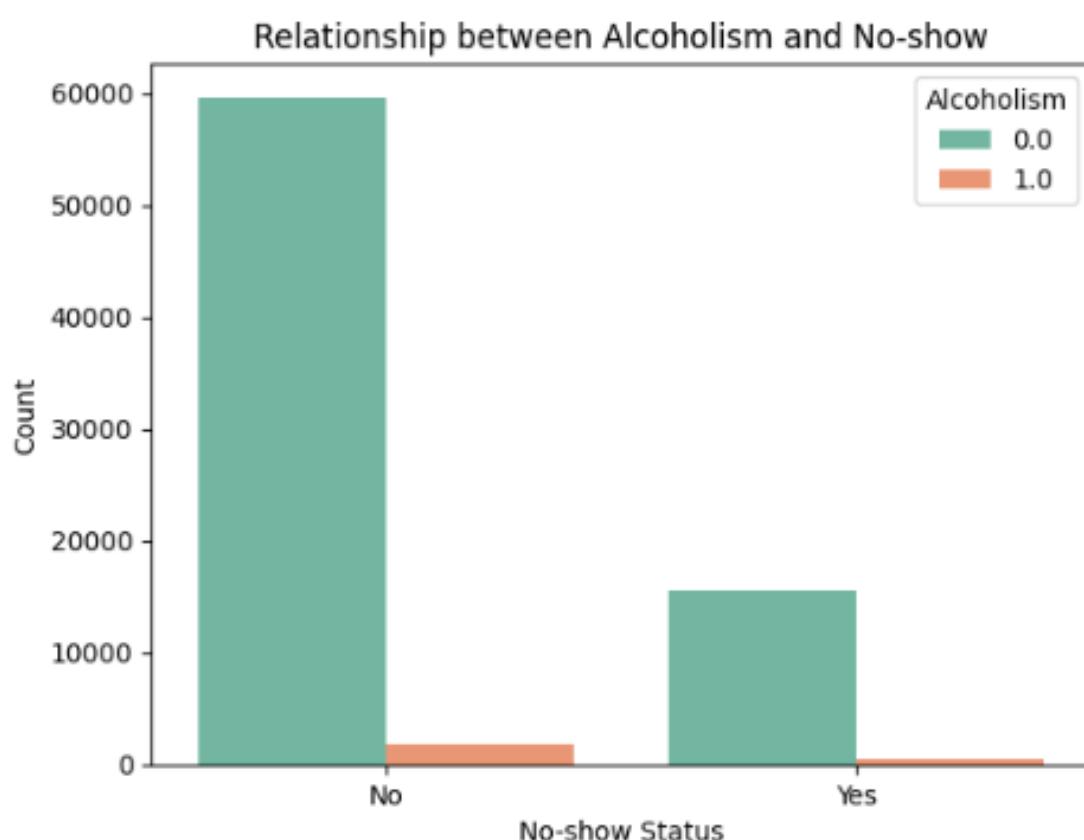
SOMAIYA
VIDYAVIHAR UNIVERSITY

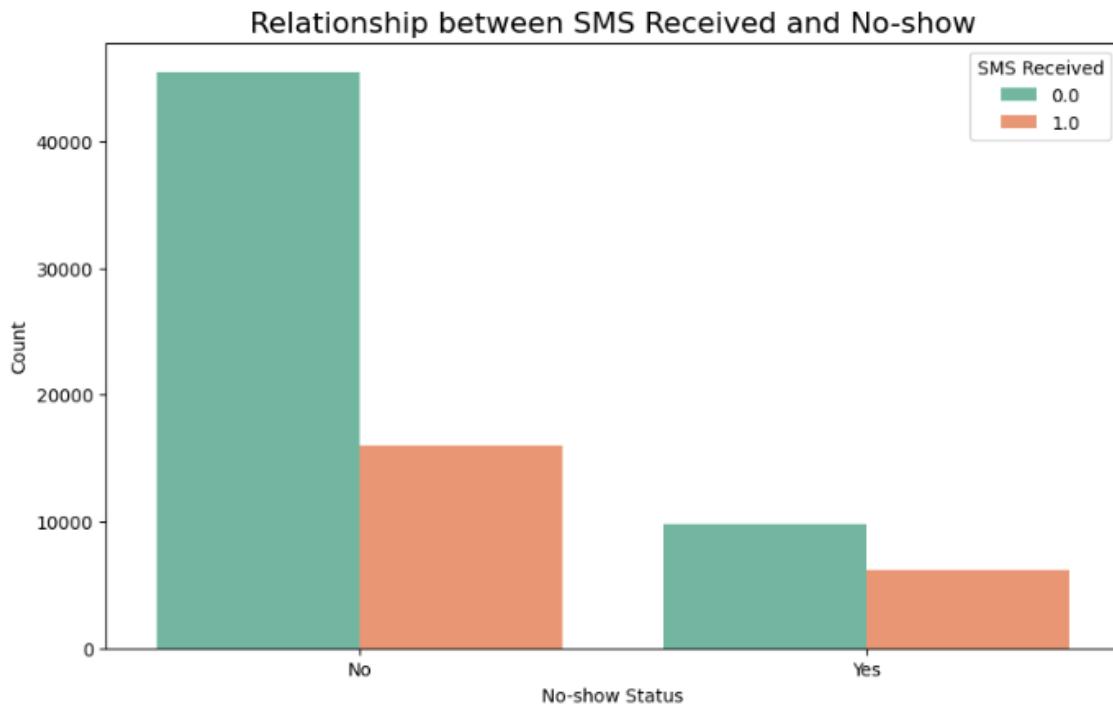
K J Somaiya School of Engineering
(formerly K J Somaiya College of Engineering)

K. J. Somaiya School of Engineering, Mumbai-77
(Somaiya Vidyavihar University)
Department of Computer Engineering

Somaiya
T R U S T







Implementation details:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from prophet import Prophet
import kagglehub
from kagglehub import KaggleDatasetAdapter

file_path = "AirPassengers.csv"

df_local = kagglehub.dataset_load(
    KaggleDatasetAdapter.PANDAS,
    "ashfakyeafi/air-passenger-data-for-time-series-analysis",
    file_path
)

df_local = df_local.rename(columns={'#Passengers': 'Passengers'})

attendance_daily = df_local[['Month', 'Passengers']].copy()
attendance_daily.columns = ['ds', 'y']

attendance_daily['ds'] = pd.to_datetime(attendance_daily['ds']).dt.tz_localize(None)

attendance_daily['y'] = pd.to_numeric(attendance_daily['y'], errors='coerce')

attendance_daily = attendance_daily.dropna()
```



```
model = Prophet()  
  
model.fit(attendance_daily)  
  
future_dates = model.make_future_dataframe(periods=30, freq='MS')  
  
forecast = model.predict(future_dates)  
  
fig = model.plot(forecast)  
plt.title('Forecast of Monthly Passenger Attendance', fontsize=16)  
plt.xlabel('Date')  
plt.ylabel('Attendance')  
plt.show()  
  
fig2 = model.plot_components(forecast)  
plt.show()  
  
def plot_passenger_time_series():  
    """Plot monthly passenger count over time"""  
    plt.figure(figsize=(10, 6))  
    sns.lineplot(x='ds', y='y', data=attendance_daily)  
    plt.title('Monthly Passenger Count Over Time', fontsize=16)  
    plt.xlabel('Date')  
    plt.ylabel('Passenger Count')  
    plt.show()  
  
plot_passenger_time_series()
```

Output:

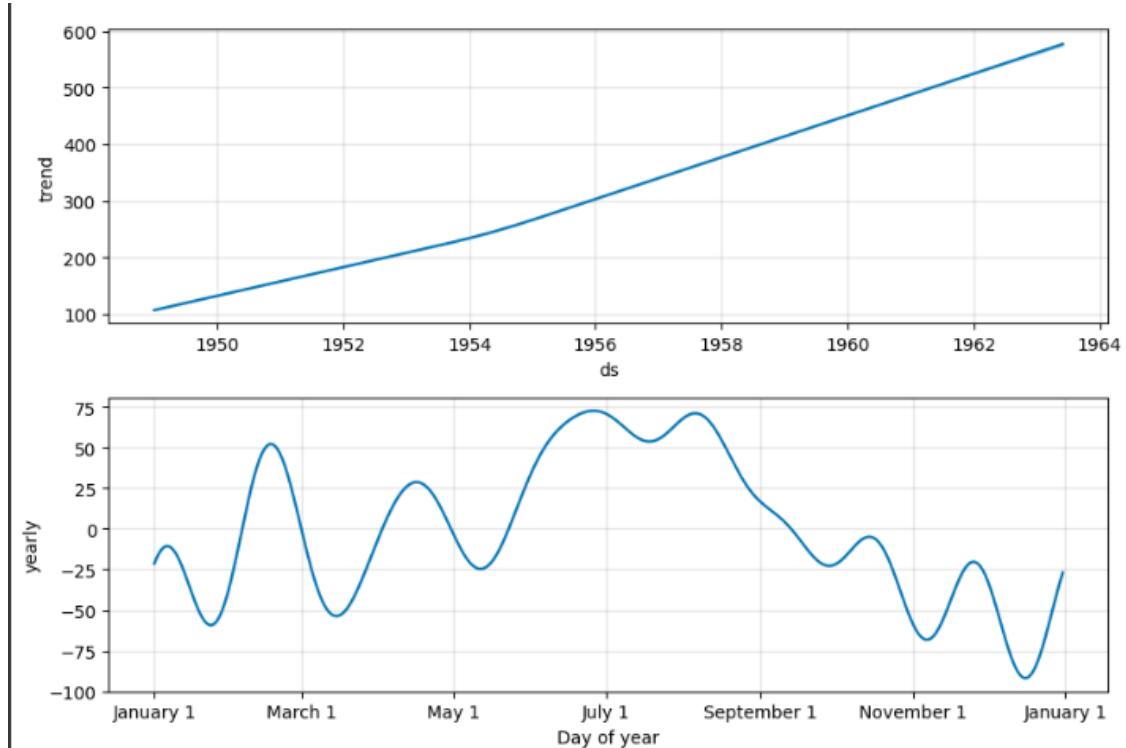
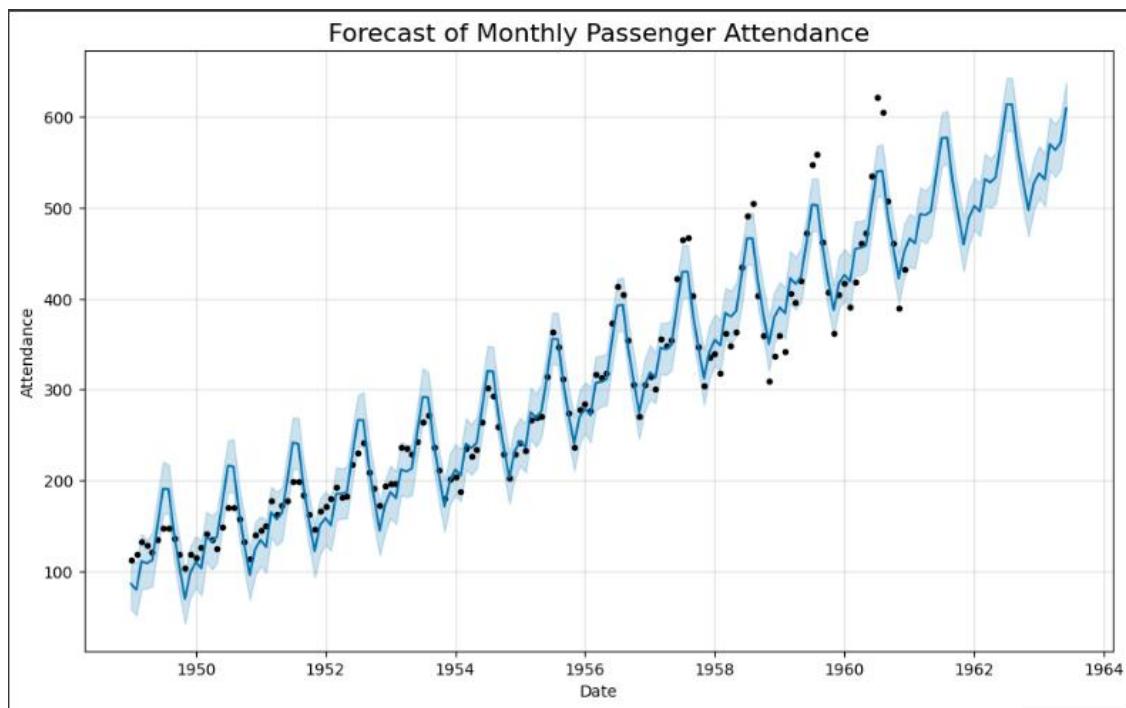


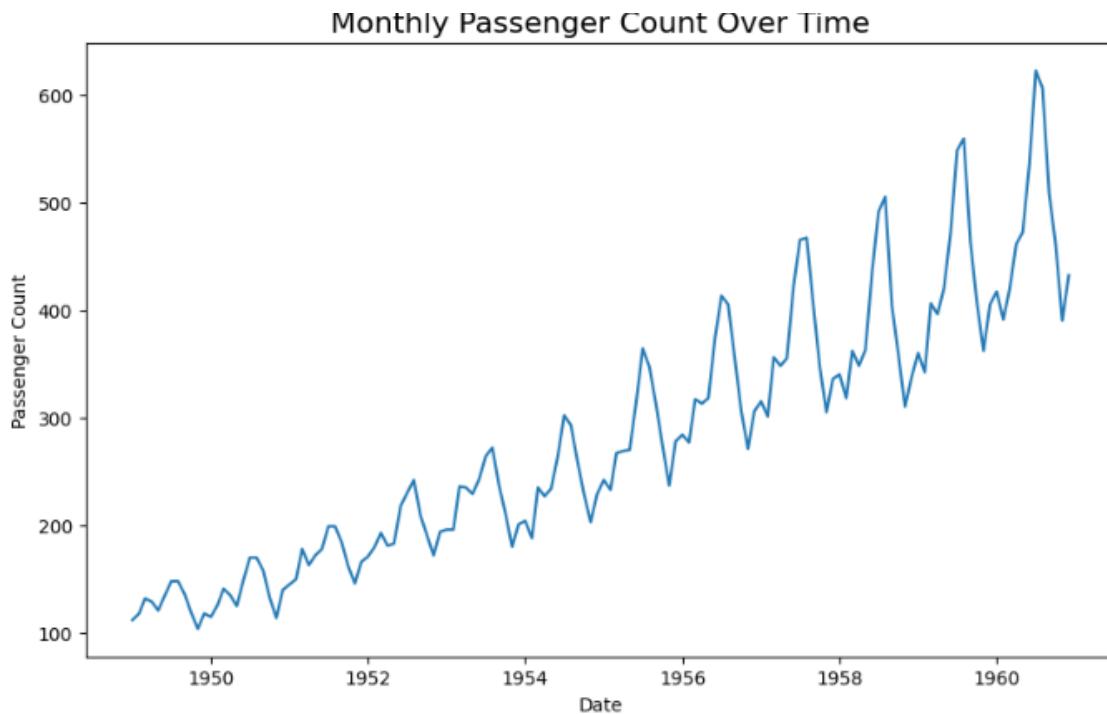
SOMAIYA
VIDYAVIHAR UNIVERSITY

K J Somaiya School of Engineering
(formerly K J Somaiya College of Engineering)

K. J. Somaiya School of Engineering, Mumbai-77
(Somaiya Vidyavihar University)
Department of Computer Engineering

Somaiya
T R U S T





Date: 13/10/2025

Signature of faculty in-charge

Post Lab Descriptive Questions:

1. Explain the components of time series?

A time series is made up of four main components: the trend, which shows the overall long-term direction of the data; seasonality, which reflects regular and repeating patterns over fixed intervals like months or quarters; cyclic patterns, which are longer-term fluctuations influenced by economic or business cycles but are irregular in length; and irregular or noise, which represents random, unpredictable variations in the data.

2. How do you handle seasonality in time series data? What methods or transformations can you apply?

Seasonality can be managed by decomposing the series to isolate and remove the seasonal component using techniques like STL decomposition or classical additive/multiplicative methods. Another approach is seasonal differencing, where values from the same season in previous cycles are subtracted to stabilize the data. You can also use seasonal dummy variables in models or apply seasonal ARIMA (SARIMA) models that explicitly account for seasonality.

3. What are some common metrics for evaluating forecasting models (e.g., MAE, RMSE, and MAPE)?

To evaluate forecasting accuracy, metrics like MAE (Mean Absolute Error) measure the average magnitude of errors without considering direction. RMSE



SOMAIYA
VIDYAVIHAR UNIVERSITY

K J Somaiya School of Engineering
(formerly K J Somaiya College of Engineering)

K. J. Somaiya School of Engineering, Mumbai-77
(Somaiya Vidyavihar University)
Department of Computer Engineering

Somaiya
T R U S T

(Root Mean Squared Error) penalizes larger errors more strongly, making it sensitive to outliers. MAPE (Mean Absolute Percentage Error) expresses errors as percentages, which is intuitive but can be misleading when actual values are near zero.